Introduction to SQL

Outline

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory

- Renamed Structured Query Language (SQL)

- ANSI and ISO standard SQL:
  - SQL-86
  - SQL-89
  - SQL-92
  - SQL:1999 (language name became Y2K compliant!)
  - SQL:2003

- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
  - Not all examples here may work on your particular system.

Data Definition Language

The SQL data-definition language (DDL) allows the specification of information about relations, including:

- The schema for each relation.

- The domain of values associated with each attribute.

- Integrity constraints

- And as we will see later, also other information such as

  - The set of indices to be maintained for each relations.

  - Security and authorization information for each relation.

  - The physical storage structure of each relation on disk.

Domain Types in SQL

- **char(n).** Fixed length character string, with user-specified length *n.*
- **varchar(n).** Variable length character strings, with user-specified maximum length *n.*
- **int.** Integer (a finite subset of the integers that is machine-dependent).
- **smallint.** Small integer (a machine-dependent subset of the integer domain type).
- **numeric(p,d).** Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (ex., **numeric**(3,1), allows 44.5 to be stores exactly, but not 444.5 or 0.32)
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(n).** Floating point number, with user-specified precision of at least *n* digits.
- More are covered in Chapter 4.

Create Table Construct

- An SQL relation is defined using the **create table** command:

    create table r (A_1 D_1, A_2 D_2, ..., A_n D_n,
            (integrity-constraint$_1$),
            ...,
            (integrity-constraint$_k$))

    - r is the name of the relation

    - each A_i is an attribute name in the schema of relation r

    - D_i is the data type of values in the domain of attribute A_i

- Example:

    create table *instructor* (
        ID   **char**(5),
        name   **varchar**(20),
        dept_name **varchar**(20),
        salary   **numeric**(8,2))

Integrity Constraints in Create Table

- **not null**

- **primary key** $(A_1, ..., A_n)$

- **foreign key** $(A_m, ..., A_n)$ **references** *r*

Example:

> **create table** *instructor* (
>   *ID*   **char**(5),
>   *name*   **varchar**(20) **not null,**
>   *dept_name* **varchar**(20),
>   *salary*   **numeric**(8,2),
>   **primary key** (*ID*),
>   **foreign key** *(dept_name)* **references** *department);*

primary key declaration on an attribute automatically ensures **not null**

And a Few More Relation Definitions

- **create table** *student* (
      ID   **varchar**(5),
      name   **varchar**(20) not null,
      dept_name   **varchar**(20),
      tot_cred   **numeric**(3,0),
      primary key *(ID),*
      foreign key *(dept_name)* **references** *department*);

- **create table** *takes* (
      ID   **varchar**(5),
      course_id   **varchar**(8),
      sec_id   **varchar**(8),
      semester   **varchar**(6),
      year   **numeric**(4,0),
      grade   **varchar**(2),
      primary key *(ID, course_id, sec_id, semester, year)* ,
      foreign key *(ID)* **references** *student,*
      foreign key *(course_id, sec_id, semester, year)* **references** *section*);
  - Note: *sec_id* can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester

And more still

- **create table** *course* (

course_id	**varchar**(8),
title	**varchar(**50),
dept_name	**varchar**(20),
credits	**numeric**(2,0),

  primary key *(course_id),*
  foreign key *(dept_name)* **references** *department);*

Updates to tables

- **Insert**
  - **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);
- **Delete**
  - Remove all tuples from the *student* relation
    - **delete from** *student*
- **Drop Table**
  - **drop table** *r*
- **Alter**
  - **alter table** *r* **add** *A D*
    - where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
    - All exiting tuples in the relation are assigned *null* as the value for the new attribute.
  - **alter table** *r* **drop** *A*
    - where *A* is the name of an attribute of relation *r*
    - Dropping of attributes not supported by many databases.

Basic Query Structure

- A typical SQL query has the form:

$$\textbf{select } A_1, A_2, ..., A_n$$
$$\textbf{from } r_1, r_2, ..., r_m$$
$$\textbf{where } P$$

- A_i represents an attribute
- R_i represents a relation
- P is a predicate.

- The result of an SQL query is a relation.

The select Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:
  > **select** *name*
  > **from** *instructor*
- NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
  - E.g., *Name* ≡ *NAME* ≡ *name*
  - Some people use upper case wherever we use bold font.

The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.

- To force the elimination of duplicates, insert the keyword **distinct** after select**.**

- Find the department names of all instructors, and remove duplicates

  select distinct *dept_name*
  from *instructor*

- The keyword **all** specifies that duplicates should not be removed.

  select all *dept_name*
  from *instructor*

The select Clause (Cont.)

- An asterisk in the select clause denotes "all attributes"

  select *
  from *instructor*

- An attribute can be a literal with no **from** clause

  select '437'

  - Results is a table with one column and a single row with value "437"
  - Can give the column a name using:

  select '437' **as** *FOO*

- An attribute can be a literal with **from** clause

  select 'A'
  from *instructor*

  - Result is a table with one column and *N* rows (number of tuples in the *instructors* table), each row with value "A"

The select Clause (Cont.)

■ The **select** clause can contain arithmetic expressions involving the operation, +, −, ∗, and /, and operating on constants or attributes of tuples.

- The query:

        select *ID, name, salary/12*
        from *instructor*

    would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename "s*alary/12*" using the **as** clause:

        select *ID, name, salary/12* **as** *monthly_salary*

The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

  select *name*
  from *instructor*
  where *dept_name* = 'Comp. Sci.'

- Comparison results can be combined using the logical connectives **and, or,** and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000

    select *name*
    from *instructor*
    where *dept_name* = 'Comp. Sci.' **and** *salary* > 80000

- Comparisons can be applied to results of arithmetic expressions.

The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor X teaches*

  select ∗
  from *instructor, teaches*

  - generates every possible instructor – teaches pair, with all attributes from both relations.

  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).

Cartesian Product

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Pinance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Pinance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Pinance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...

Examples

- Find the names of all instructors who have taught some course and the course_id

  - **select** *name, course_id*
    from *instructor , teaches*
    where *instructor.ID = teaches.ID*

- Find the names of all instructors in the Art department who have taught some course and the course_id

  - **select** *name, course_id*
    from *instructor , teaches*
    where *instructor.ID = teaches.ID* **and** *instructor. dept_name* = 'Art'

The Rename Operation

■ The SQL allows renaming relations and attributes using the **as** clause:

old-name **as** *new-name*

■ Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

- **select distinct** *T.name*
  from *instructor* **as** *T, instructor* **as** *S*
  where *T.salary > S.salary* **and** *S.dept_name = 'Comp. Sci.'*

■ Keyword **as** is optional and may be omitted
instructor **as** *T ≡ instructor T*

Self Join Example

- Relation *emp-super*

person	*supervisor*
Bob	Alice
Mary	Susan
Alice	David
David	Mary

- Find the supervisor of "Bob"

- Find the supervisor of the supervisor of "Bob"

- Find ALL the supervisors (direct and indirect) of "Bob

String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:

  - percent (%). The % character matches any substring.

  - underscore (_). The _ character matches any character.

- Find the names of all instructors whose name includes the substring "dar".

  > **se**l**ect** *name*
  > **from** *instructor*
  > **where** *name* **like** '%dar%'

- Match the string "100%"

  > **like** '100 \%' **escape** '\'

  in that above we use backslash (\) as the escape character.

String Operations (Cont.)

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro".
  - '%Comp%' matches any string containing "Comp" as a substring.
  - '_ _ _' matches any string of exactly three characters.
  - '_ _ _ %' matches any string of at least three characters.

- SQL supports a variety of string operations such as
  - concatenation (using "||")
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

    select distinct *name*
    from   *instructor*
    order by *name*

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

    - Example: **order by** *name* **desc**

- Can sort on multiple attributes

    - Example: **order by** *dept_name, name*

Where Clause Predicates

- SQL includes a **between** comparison operator

- Example: Find the names of all instructors with salary between $90,000 and $100,000 (that is, $\geq$ $90,000 and $\leq$ $100,000)

  - **select** *name*
    from *instructor*
    where *salary* **between** 90000 **and** 100000

- Tuple comparison

  - **select** *name*, *course_id*
    from *instructor*, *teaches*
    where (*instructor*.*ID*, *dept_name*) = (*teaches*.*ID*, 'Biology');

Duplicates

- In relations with duplicates, SQL can define how many copies of tuples appear in the result.

- **Multiset** versions of some of the relational algebra operators – given multiset relations r_1 and r_2:

  1. $\sigma_\theta(r_1)$: If there are c_1 copies of tuple t_1 in r_1, and t_1 satisfies selections σ_θ, then there are c_1 copies of t_1 in $\sigma_\theta(r_1)$.

  2. $\Pi_A(r)$: For each copy of tuple t_1 in r_1, there is a copy of tuple $\Pi_A(t_1)$ in $\Pi_A(r_1)$ where $\Pi_A(t_1)$ denotes the projection of the single tuple t_1.

  3. $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2, there are $c_1 \times c_2$ copies of the tuple $t_1. t_2$ in $r_1 \times r_2$

Duplicates (Cont.)

- Example: Suppose multiset relations r_1 (A, B) and r_2 (C) are as follows:

$$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then $\Pi_B(r_1)$ would be {(a), (a)}, while $\Pi_B(r_1)$ x r_2 would be

$$\{(a,2), (a,2), (a,3), (a,3), (a,3), (a,3)\}$$

- SQL duplicate semantics:

  select $A_1, A_2, ..., A_n$
  from $r_1, r_2, ..., r_m$
  where P

  is equivalent to the *multiset* version of the expression:

$$\prod_{A_1,A_2,...,A_n} (\sigma_P(r_1 \times r_2 \times \ldots \times r_m))$$

Set Operations

■ Find courses that ran in Fall 2009 or in Spring 2010

(**select** *course_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2009)
 union
(**select** *course_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2010)

■ Find courses that ran in Fall 2009 and in Spring 2010

(**select** *course_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2009)
 intersect
(**select** *course_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2010)

■ Find courses that ran in Fall 2009 but not in Spring 2010

(**select** *course_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2009)
 except
(**select** *course_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2010)

Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.

  - **select distinct** *T.salary*
    from *instructor* **as** *T, instructor* **as** *S*
    where *T.salary < S.salary*

- Find all the salaries of all instructors

  - **select distinct** *salary*
    from *instructor*

- Find the largest salary of all instructors.

  - (**select** "second query")
    except
     (**select** "first query")

Set Operations (Cont.)

- Set operations **union, intersect,** and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all, intersect all** and **except all.**

- Suppose a tuple occurs m times in r and n times in $s,$ then, it occurs:
  - $m + n$ times in r **union all** s
  - $\min(m,n)$ times in r **intersect all** s
  - $\max(0, m - n)$ times in r **except all** s

Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes

- *null* signifies an unknown value or that a value does not exist.

- The result of any arithmetic expression involving *null* is *null*

  - Example: 5 + *null* returns null

- The predicate **is null** can be used to check for null values.

  - Example: Find all instructors whose salary is null*.*

    select *name*
    from *instructor*
    where *salary* **is null**

Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*

- Any comparison with *null* returns *unknown*
  - Example*: 5 < null   or   null <> null   or   null = null*

- Three-valued logic using the value *unknown*:
  - OR: (*unknown* **or** *true*)   = *true*,
        (*unknown* **or** *false*) = *unknown*
        (*unknown* **or** *unknown*) = *unknown*

  - AND: *(true* **and** *unknown) = unknown,*
         (false **and** *unknown) = false,*
         (unknown **and** *unknown) = unknown*

  - NOT*: (***not** *unknown) = unknown*

  - "*P* **is unknown**" evaluates to true if predicate *P* evaluates to *unknown*

- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

  avg: average value
  min: minimum value
  max: maximum value
  sum: sum of values
  count: number of values

Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department

  - **select avg** (*salary*)
    from *instructor*
    where *dept_name*= 'Comp. Sci.';

- Find the total number of instructors who teach a course in the Spring 2010 semester

  - **select count** (**distinct** *ID*)
    from *teaches*
    where *semester* = 'Spring' **and** *year* = 2010;

- Find the number of tuples in the *course* relation

  - **select count** (*)
    from *course*;

Aggregate Functions – Group By

- Find the average salary of instructors in each department

    - **select** *dept_name*, **avg** (*salary*) **as** *avg_salary*
      from *instructor*
      group by *dept_name*;

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
    - /* erroneous query */
      select *dept_name*, *ID*, **avg** (*salary*)
      from *instructor*
      group by *dept_name*;

Aggregate Functions – Having Clause

■ Find the names and average salaries of all departments whose average salary is greater than 42000

> **select** *dept_name*, **avg** (*salary*)
> **from** *instructor*
> **group by** *dept_name*
> **having avg** (*salary*) > 42000;

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

Null Values and Aggregates

- Total all salaries

    > **select sum** (*salary*)
    > **from** *instructor*

    - Above statement ignores null amounts
    - Result is *null* if there is no non-null amount

- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes

- What if collection has only null values?

    - count returns 0
    - all other aggregates return null

Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.

- The nesting can be done in the following SQL query

    select $A_1, A_2, ..., A_n$
    from $r_1, r_2, ..., r_m$
    where P

  as follows:

    - A_i can be replaced be a subquery that generates a single value.

    - r_i can be replaced by any valid subquery

    - P can be replaced with an expression of the form:

            B <operation> (subquery)

    Where B is an attribute and <operation> to be defined later.

Subqueries in the Where Clause

Subqueries in the Where Clause

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality.

Set Membership

- Find courses offered in Fall 2009 and in Spring 2010

  select distinct *course_id*
  from *section*
  where *semester* = 'Fall' **and** *year*= 2009 **and**
  　　course_id **in** (**select** *course_id*
  　　　　　　　from *section*
  　　　　　　　where *semester* = 'Spring' **and** *year*= 2010);

- Find courses offered in Fall 2009 but not in Spring 2010

  select distinct *course_id*
  from *section*
  where *semester* = 'Fall' **and** *year*= 2009 **and**
  　　course_id **not in** (**select** *course_id*
  　　　　　　　from *section*
  　　　　　　　where *semester* = 'Spring' **and** *year*= 2010);

Set Membership (Cont.)

- Find the total number of (distinct) students who have taken course sections taught by the instructor with *ID* 10101

   select count (**distinct** *ID*)
   from *takes*
   where (*course_id*, *sec_id*, *semester*, *year*) **in**
             (**select** *course_id*, *sec_id*, *semester*, *year*
             from *teaches*
             where *teaches.ID*= 10101);

- Note: Above query can be written in a much simpler manner.
  The formulation above is simply to illustrate SQL features.

Set Comparison – "some" Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

  select distinct *T.name*
  from *instructor* **as** *T, instructor* **as** *S*
  where *T.salary* > *S.salary* **and** *S.dept name* = 'Biology';

- Same query using > **some** clause

  select *name*
  from *instructor*
  where *salary* > **some** (**select** *salary*
                       from *instructor*
                       where *dept name* = 'Biology');

Definition of "some" Clause

- F <comp> **some** $r \Leftrightarrow \exists\, t \in r$ such that (F <comp> t)
  Where <comp> can be: $<, \leq, >, =, \neq$

(5 < **some** $\begin{array}{|c|}\hline 0 \\\hline 5 \\\hline 6 \\\hline\end{array}$) = true

(read: 5 < some tuple in the relation)

(5 < **some** $\begin{array}{|c|}\hline 0 \\\hline 5 \\\hline\end{array}$) = false

(5 = **some** $\begin{array}{|c|}\hline 0 \\\hline 5 \\\hline\end{array}$) = true

(5 ≠ **some** $\begin{array}{|c|}\hline 0 \\\hline 5 \\\hline\end{array}$) = true (since 0 ≠ 5)

(= **some**) ≡ **in**
However, (≠ **some**) $\not\equiv$ **not in**

Set Comparison – "all" Clause

■ Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

select *name*
from *instructor*
where *salary* > **all** (**select** *salary*
                          from *instructor*
                          where *dept name* = 'Biology');

Definition of "all" Clause

■ F <comp> **all** $r \Leftrightarrow \forall\, t \in r$ (F <comp> t)

(5 < **all**
0
5
6
) = false

(5 < **all**
6
10
) = true

(5 = **all**
4
5
) = false

(5 ≠ **all**
4
6
) = true (since 5 ≠ 4 and 5 ≠ 6)

(≠ **all**) ≡ **not in**
However, (= **all**) ≢ **in**

Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.

- **exists** $r \Leftrightarrow r \neq \emptyset$

- **not exists** $r \Leftrightarrow r = \emptyset$

Use of "exists" Clause

- Yet another way of specifying the query "Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester"

  select *course_id*
  from *section* **as** *S*
  where *semester* = 'Fall' **and** *year* = 2009 **and**
        exists (**select** *
              from *section* **as** *T*
              where *semester* = 'Spring' **and** *year*= 2010
                    and *S.course_id* = *T.course_id*);

- **Correlation name** – variable S in the outer query
- **Correlated subquery** – the inner query

Use of "not exists" Clause

- Find all students who have taken all courses offered in the Biology department.

  ```
  select distinct S.ID, S.name
  from student as S
  where not exists ( (select course_id
                         from course
                         where dept_name = 'Biology')
                     except
                       (select T.course_id
                         from takes as T
                         where S.ID = T.ID));
  ```

  - First nested query lists all courses offered in Biology
  - Second nested query lists all courses a particular student took

- Note that $X - Y = \emptyset \iff X \subseteq Y$

- *Note:* Cannot write this query using = **all** and its variants

Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.

- The **unique** construct evaluates to "true" if a given subquery contains no duplicates .

- Find all courses that were offered at most once in 2009

    select *T.course_id*
    from *course* **as** *T*
    where unique (**select** *R.course_id*
                            from *section* **as** *R*
                            where *T.course_id*= *R.course_id*
                                and *R.year* = 2009);

Subqueries in the Form Clause

Subqueries in the Form Clause

- SQL allows a subquery expression to be used in the **from** clause

- Find the average instructors' salaries of those departments where the average salary is greater than $42,000."

  > **select** *dept_name*, *avg_salary*
  > **from** (**select** *dept_name*, **avg** (*salary*) **as** *avg_salary*
  >   **from** *instructor*
  >   **group by** *dept_name*)
  > **where** *avg_salary* > 42000;

- Note that we do not need to use the **having** clause

- Another way to write above query

  > **select** *dept_name*, *avg_salary*
  > **from** (**select** *dept_name*, **avg** (*salary*)
  >   **from** *instructor*
  >   **group by** *dept_name*) **as** *dept_avg* (*dept_name*, *avg_salary*)
  >
  > **where** *avg_salary* > 42000;

With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.

- Find all departments with the maximum budget

   with *max_budget* (*value*) **as**
         (**select max**(*budget*)
          from *department*)
   select *department.name*
   from *department*, *max_budget*
   where *department.budget = max_budget.value*;

Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

  with *dept _total* (*dept_name*, *value*) **as**
      (**select** *dept_name*, **sum**(*salary*)
     from *instructor*
     group by *dept_name*),
  dept_total_avg(*value*) **as**
      (**select avg**(*value*)
    from *dept_total*)
  select *dept_name*
  from *dept_total*, *dept_total_avg*
  where *dept_total.value* > *dept_total_avg.value*;

Subqueries in the Select Clause

Scalar Subquery

- Scalar subquery is one which is used where a single value is expected

- List all departments along with the number of instructors in each department

select *dept_name*,
        (**select count**(*)
          from *instructor*
          where *department*.*dept_name* = *instructor*.*dept_name*)
        as *num_instructors*
from *department*;

- Runtime error if subquery returns more than one result tuple

Modification of the Database

- Deletion of tuples from a given relation.

- Insertion of new tuples into a given relation

- Updating of values in some tuples in a given relation

Deletion

■ Delete all instructors

          delete from *instructor*

■ Delete all instructors from the Finance department

          delete from *instructor*
          where *dept_name*= 'Finance';

■ Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

          delete from *instructor*
          where *dept name* **in** (**select** *dept name*
                           from *department*
                           where *building* = 'Watson');

Deletion (Cont.)

■ Delete all instructors whose salary is less than the average salary of instructors

     delete from *instructor*
     where *salary* < (**select avg** (*salary*)
                    from *instructor*);

● Problem: as we delete tuples from deposit, the average salary changes

● Solution used in SQL:

    1. First, compute **avg** (salary) and find all tuples to delete

    2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

Insertion

- Add a new tuple to *course*

  insert into *course*
      values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- or equivalently

  insert into *course* (*course_id*, *title*, *dept_name*, *credits*)
      values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- Add a new tuple to *student* with *tot_creds* set to null

  insert into *student*
      values ('3003', 'Green', 'Finance', *null*);

Insertion (Cont.)

■ Add all instructors to the *student* relation with tot_creds set to 0

> **insert into** *student*
>   **select** *ID, name, dept_name, 0*
>   **from** *instructor*

■ The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

> **insert into** *table*1 **select** * **from** *table*1

would cause problem

Updates

- Increase salaries of instructors whose salary is over $100,000 by 3%, and all others by a 5%

  - Write two **update** statements:

    update *instructor*
        set *salary = salary* * 1.03
        where *salary* > 100000;
    update *instructor*
        set *salary = salary* * 1.05
        where *salary* <= 100000;

  - The order is important

  - Can be done better using the **case** statement (next slide)

Case Statement for Conditional Updates

- Same query as before but with case statement

  update *instructor*
     set *salary* = **case**
                      when *salary* <= 100000 **then** *salary* * 1.05
                      else *salary* * 1.03
                      end

Updates with Scalar Subqueries

- Recompute and update tot_creds value for all students

  update *student S*
      set *tot_cred* = (**select sum**(*credits*)
                             from *takes, course*
                             where *takes.course_id = course.course_id* **and**
                                      *S.ID= takes.ID.***and**
                                      takes.grade <> 'F' **and**
                                      takes.grade **is not null**);

- Sets *tot_creds* to null for students who have not taken any course

- Instead of **sum**(*credits*), use:

  case
      when sum(*credits*) **is not null then sum**(*credits*)
      else 0
  end